WACV #99

Unlearn in Context: Contrastive Alignment via Diversified Supervision for Diffusion Models

Anonymous WACV Algorithms Track submission

Paper ID 99

Abstract

Concept unlearning has emerged as a promising direction for reducing the risks of harmful content generation in text-to-image diffusion models by selectively erasing undesirable concepts from a model's parameters. Existing approaches typically rely on keywords to identify the target concept to be unlearned. However, we show that this keyword-based formulation is inherently limited: a visual concept is multi-dimensional, can be expressed in diverse textual forms, and often overlaps with related concepts in the latent space. We therefore adopt a contextual view: a concept should be unlearned through the company it keeps. Under this view, keyword-based unlearning approaches, which imprecisely indicate the target concept via a fixed representation, are brittle against prompt-rephrasing jailbreaks and highly prone to collateral over-forgetting. To address these limitations, we propose ***Diversified Unlearning****, a framework that exposes the erasure objective to natural context-dependent embedding variations. Rather than relying on simple prompt augmentation, we construct paired contextual formulations that introduce* ***implicit contrastive supervision*** *directly into the forgetting process, encouraging the model to isolate the target concept while keeping surrounding contexts invariant. Through extensive experiments across multiple benchmarks and state-of-the-art baselines, we demonstrate that integrating our framework as a lightweight add-on into existing unlearning pipelines yields notable improvements in both erasure effectiveness and concept retention, while significantly hardening models against adversarial recovery attacks. All experimental results and detailed implementations can be found at* `https://anonymous.4open.science/r/Diversified_Unlearning`.

1. Introduction

Recent text-to-image generative models, such as Stable Diffusion [31] or Dall-E [29], inadvertently learn **undesirable concepts** including *toxicity, social biases, and copyrighted material* from web-scale data [33]. Consequently, these models can be exploited to generate harmful or infringing content [6, 37], posing severe societal risks. While data filtering before training (pre-processing) and output censorship (post-processing) mitigate these risks, both remain limited by high computational costs and inefficiency under continuous large-scale queries.

A more promising direction is ***concept unlearning***, which directly removes undesired concepts from a model by fine-tuning its parameters [4, 5, 10, 11, 22, 40]. Existing approaches generally fall into two categories: *Output-based methods* [5, 10] that force the noise prediction of a target concept c_e to match a neutral anchor c_a, and *Attention-based methods* [11, 22] that modify cross-attention weights to weaken text-to-visual alignments.

Despite their intuitive appeal, both families are typically formulated in a ***keyword-based*** manner, representing a visual concept by isolated textual tokens (e.g., "Barack Obama" or "a photo of Barack Obama"). However, visual concepts are not tied to a single lexical form: the same target can appear through paraphrases, attributes, actions, and surrounding scenes. This limitation echoes a classic distributional principle in linguistics: "You shall know a word by the company it keeps" [9]. We adapt this principle to concept unlearning: *one shall unlearn a word by the company it keeps*. In text-to-image models, this "company" consists of the natural contexts in which a concept appears, including its attributes, actions, co-occurring objects, backgrounds, and paraphrased textual descriptions. Reliable unlearning should therefore identify the target across these contextual variations while preserving the surrounding context itself. Through systematic analysis, we identify two failure modes showing why keyword-based unlearning falls short of this contextual principle:

- **Observation 1: Imprecise Target Representation.** During pre-training, concepts are entangled with diverse linguistic contexts, forming context-dependent representation regions rather than isolated points. Our perturbation

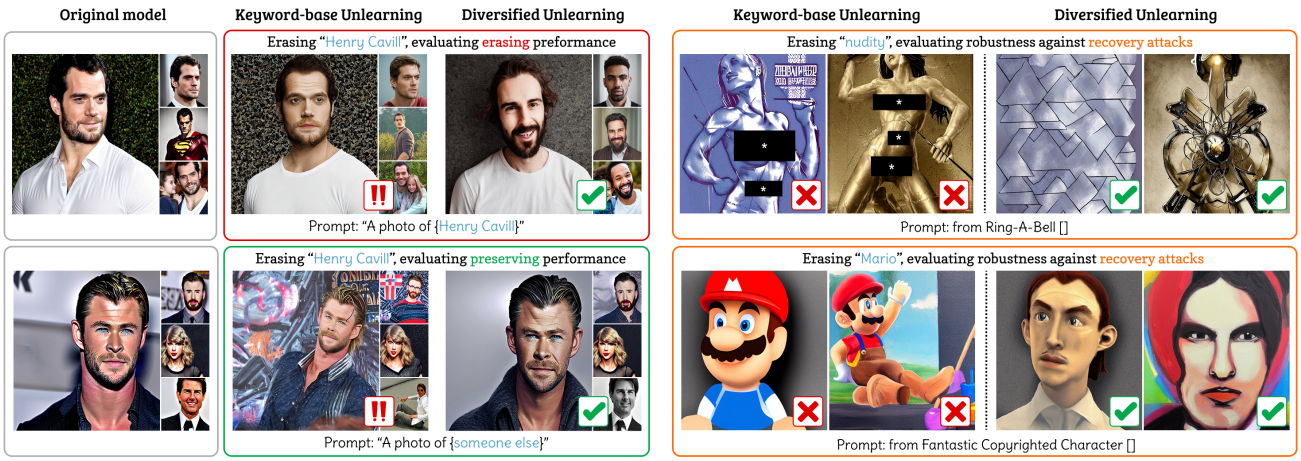

Figure 1. Diversified Unlearning enhances the representation of target concepts on top of state-of-the-art methods, significantly improves both erasing performance on target concepts and preserving abilities on unrelated ones. Moreover, our method also demonstrated effectiveness in mitigating recovery attacks against sexual concepts and copyrighted characters. (✓ good performance; ✗ bad performance; !! moderate performance). (*) Censoring for publication.

analysis (Table 1) shows that, even after adding noise to the text embedding, the underlying diffusion model can still generate the target concept c_e with high probability. This indicates that the concept remains recoverable from a neighborhood of semantically related embeddings. Keyword-based unlearning, which uses a single fixed representation to identify the target, may therefore suppress the canonical keyword while leaving nearby contextual variants untouched, making the model susceptible to **prompt rephrasing jailbreaks** and noise-based recovery attacks (Table 5).

- **Observation 2: Ambiguous Erasure Signals in Output Space.** In output-based unlearning, a prompt containing an isolated keyword c_e is not decoded in isolation: the stochastic reverse process often also produces co-occurring objects, attributes, and background contexts in intermediate latents. Forcing the full prediction trajectory to match a generic anchor c_a provides an ambiguous training signal because the loss does not specify which visual factors correspond to the target and which should be preserved. As a result, benign contextual features can be suppressed together with the target. This issue is further exacerbated by traditional data augmentation or adversarial unlearning strategies [1, 17, 47]; when diversified or adversarially perturbed inputs are mapped to an unconstrained anchor, they can introduce an **overly aggressive erasure signal** that compromises preservation (Table 2).

Motivated by this contextual view, we propose **Diversified Unlearning**, a framework that unlearns a concept through the company it keeps: diverse contextual variations whose non-target content is structurally controlled. This design introduces **Implicit Contrastive Supervision**: the model is encouraged to change the target concept while preserving the context shared by the target and anchor pair.

For output-based methods, we introduce **Diversified Prompting**. Instead of mapping an isolated keyword c_e to a generic anchor c_a, we construct paired formulations: $\{c + c_e\} \rightarrow \{c + c_a\}$, where $c \sim \mathcal{C}$ denotes clean, natural contexts. Because the context c is shared across both sides of the transformation, the erasure direction is concentrated on the difference between c_e and c_a. This context-invariant formulation helps the model distinguish *what to erase* from *what to preserve*, reducing collateral damage to adjacent concepts and background contexts.

For attention-based methods, direct context manipulation can cause severe over-forgetting due to text encoder semantic bias. Because the encoder aggregates semantics from both the target and the context, suppressing the full prompt embedding indiscriminately erases *shared contextual features*, harming general generative ability. We thus propose **Diversified Embedding Mixup**, a ***token-level interpolation*** strategy that blends target tokens with contextualized embeddings. Crucially, these blended contexts are explicitly incorporated into the preservation objective, encouraging the model to isolate the target identity while safeguarding surrounding contextual features against representation collapse.

As lightweight add-ons, our modules enhance existing unlearning methods by providing richer textual coverage while mitigating the instability and over-forgetting issues. Evaluation across five diverse settings–celebrities, physical objects, copyrighted characters, nudity, and artistic styles (Table 1, Supplementary Material)–demonstrates substantial improvements in both erasure and preservation. For instance, in celebrity erasure (Table 3), our framework achieves a **∼50% reduction** in GPT-score [25]. More importantly, our approach exhibits strong resilience against advanced adversarial recovery attacks, including Ring-A-

Bell [39], Fantastic Copyrighted Concepts [14], and Noise-based Attacks [21]. This structural **robustness**, alongside improved concept **erasure** and **preservation**, manifests as the **triple benefits** highlighted in Figure 1, where upgrading baselines with our method suppresses the regeneration of sensitive targets like "nudity" and "Mario".

In summary, our main contributions are:

① We identify two key limitations of keyword-based concept unlearning: isolated keywords provide imprecise target representations, and output-space objectives can produce ambiguous erasure signals that compromise preservation. We propose **Diversified Unlearning**, a context-structured framework inspired by the principle that a concept should be unlearned through the company it keeps, introducing *implicit contrastive supervision* through paired contextual prompts and token-level embedding mixup.

② Through extensive experiments across diverse evaluation settings, we demonstrate notable improvements in **both erasure effectiveness and concept preservation**. Under adversarial recovery attacks, Diversified Unlearning also exhibits **significantly stronger robustness**, suggesting a valuable perspective on how visual concepts should be represented within the unlearning pipeline.

2. Preliminaries and Related Work

Latent Diffusion Models (LDMs). Text-conditioned LDMs [16, 29–31] generate images in a latent space compressed by a pretrained encoder $z_0 = \mathcal{E}(x)$. Given a prompt c encoded as $\tau(c)$ by a text encoder such as CLIP [28], the denoising network ϵ_θ is trained to predict noise $\epsilon \sim \mathcal{N}(0, \mathbf{I})$ at timestep t:

$$\min_{\theta} \mathbb{E}_{z_0,t,\epsilon} \left\| \epsilon - \epsilon_\theta(z_t, t, \tau(c)) \right\|_2^2 . \tag{1}$$

Concept Unlearning. Concept erasing aims to update model weights θ to a sanitized state θ' that suppresses undesirable concepts $c_e \in \mathbf{E}$ while preserving benign concepts $c_p \in \mathbf{P}$. Existing safety interventions can be grouped into inference-time and finetuning-based methods. Inference-time approaches do not permanently modify model weights: they include post-hoc filtering with detectors such as Nudenet [27], prompt rewriting or negative prompting [14, 35], and latent/attention interventions such as SLD [33], TRCE [32], SAFREE [44]. These methods are flexible, but remain vulnerable to adversarial prompting [3, 43], can be bypassed in open deployments [36], and must be applied at inference time.

Finetuning-based unlearning instead modifies the model parameters and redistributes a sanitized checkpoint. Dataset-level filtering before pre-training [34, 35, 38] is expensive and can leave residual unsafe concepts [10]; parameter-level finetuning is therefore the dominant practical alternative. Current methods broadly follow two paradigms [5]:

1) Output-based Unlearning: Methods such as ESD [10], SalUn [8], AP [4], AGE [5], and ACE [40] steer the noise prediction of an erased concept c_e toward a neutral anchor c_a, while using a preservation loss to retain benign concepts c_p:

$$\min_{\theta'} \mathbb{E}_{c_e,c_p} \Big[\left\| \epsilon_{\theta'}(z_t, t, \tau(c_e)) - \epsilon_\theta(z_t, t, \tau(c_a)) \right\|_2^2 + \left\| \epsilon_{\theta'}(z_t, t, \tau(c_p)) - \epsilon_\theta(z_t, t, \tau(c_p)) \right\|_2^2 \Big]. \tag{2}$$

2) Attention-based Unlearning: To avoid sampling noisy latent trajectories, methods such as UCE [11] and MACE [22] directly edit cross-attention projection weights W (e.g., W^K or W^V), weakening the alignment between target text features and visual tokens [11, 19, 24, 45]:

$$\min_{W'} \sum_{c_e \in \mathbf{E}} \left\| W' \tau(c_e) - W \tau(c_a) \right\|_2^2 + \sum_{c_p \in \mathbf{P}} \left\| W' \tau(c_p) - W \tau(c_p) \right\|_2^2 . \tag{3}$$

where W and W' denote the original and updated cross-attention weights. This formulation can admit closed-form updates in specific settings [24] and operates purely in embedding space.

Robust and Diversified Unlearning. Recent works augment keyword-based erasure with adversarial training, distribution alignment, or data diversification [12, 17, 18, 41, 47, 49], partly motivated by black-box [39, 42] and white-box recovery attacks [7, 48]. However, unconstrained diversification can produce overly aggressive erasure signals that harm preservation; our framework instead uses structured context pairs to separate target removal from context preservation without architectural overhead. Due to space constraints, we provide an extended discussion of recent advancements in the Supplementary Material A.

3. An Empirical Analysis of Keyword-based Unlearning

Keyword-based unlearning operationalizes concept erasure by mapping isolated tokens (e.g., $c_e =$ *"Barack Obama"*) to a neutral anchor c_a. However, this fixed-indicator formulation only partially reflects the contextual company through which concepts are represented in text-to-image (T2I) models, where a target may appear across paraphrases, attributes, actions, and surrounding scenes. In this section, we expand upon the two core observations from Section 1 with empirical evidence suggesting why such erasure can lead to fragile forgetting and suboptimal preservation.

Observation 1: Imprecise Target Representation. During T2I pre-training, visual concepts are learned within diverse linguistic contexts. Consequently, within the text encoder's output space, the embedding of a target concept c_e can vary with surrounding tokens and prompt structure. Keyword-based unlearning, however, optimizes parameters against an isolated keyword, using a single fixed text em-

WACV #99

Table 1. Embedding perturbation analysis under varying Gaussian noise ratios α added to the prompt *"a photo of Cassette Player/Church"*. DS-1 and DS-5 denote the top-1 and top-5 detection score of a pre-trained ResNet50 [13].

SD	*"**Cassette Player**" concept*						*"**Church**" concept*					
	$\alpha=0$	0.1	0.3	0.5	0.8	1.0	$\alpha=0$	0.1	0.3	0.5	0.8	1.0
DS-1	5	6	3	3	1	0	86	80	75	38	1	1
DS-5	97	100	74	37	8	0	100	100	91	54	4	2

bedding as the target indicator. This creates a potential mismatch with the context-rich nature of pre-training data.

To probe whether concept information is confined to the canonical text embedding, we inject Gaussian noise (ratio α) into the text embedding of *"a photo of Cassette Player/Church"*. As shown in Table 1, even under moderate noise ($\alpha = 0.3$), the model still generates the target concepts at non-trivial rates, with top-5 detection scores (DS-5) of 74% for *"Cassette Player"* and 91% for *"Church"*. This suggests that concept-bearing signals persist in neighborhoods around the canonical embedding. Accordingly, a keyword-only update may suppress the exact training indicator while leaving nearby contextual or perturbed variants insufficiently constrained. This interpretation is consistent with the robustness results in Table 5: standard AP [4] remains more vulnerable than its diversified counterpart under Ring-A-Bell [39] and Noise ATK [21], including a high residual GPT-score of 39.63 under Noise ATK compared with 14.94 after diversification. These results support the need to cover contextual variants rather than treating the keyword as a single point estimate.

Observation 2: Ambiguous Erasure Signals in Output Space. In output-based unlearning, a prompt containing an isolated keyword c_e is not decoded in isolation. During the reverse process, the generated trajectory can also include backgrounds, attributes, or objects associated with the target. For example, when erasing *"Barack Obama"*, intermediate outputs may include co-occurring elements such as a business suit or the American flag. Forcing these complete, multi-object trajectories to match a generic anchor c_a can make the training signal ambiguous: the objective does not explicitly specify which visual factors correspond to the target identity and which should remain preserved. As a result, the erasure signal may also affect surrounding attributes. This interpretation is consistent with the preservation degradation observed in Table 3; for instance, ESD [10] and AP [4] retain relatively low context-preservation GPT scores [25] of 31.68 and 14.97, respectively.

This issue can become more pronounced when prompt augmentation or adversarial unlearning strategies diversify the input without preserving the structure of the surrounding context [1, 17, 47]. These approaches often introduce explicit context addition $\{c_{\text{context}} + c_e\}$ or adversarial embedding perturbations to broaden coverage. However, if

Table 2. Object erasure performance across prompt complexities (Supp. Material B.1.2). Erasure Success Rate (ESR) and Preservation Success Rate (PSR) are computed at Top-1 (-1) and Top-5 (-5) (higher is better). Aug-ESD/AP and Adv denote ESD [10]/AP [4] with prompt augmentation and AdvUnlearn [47], respectively.

Object	*"**Garbage Truck**" erasure*				*"**Five objects**" erasure*			
	ESR-1	ESR-5	PSR-1	PSR-5	ESR-1	ESR-5	PSR-1	PSR-5
ESD	0.968	0.934	0.652	0.831	0.891	0.754	0.470	0.634
Aug-ESD	**1.000**	**1.000**	0.566	0.729	**0.987**	**0.958**	0.350	0.484
Diversified-ESD	0.998	0.996	**0.775**	**0.933**	0.954	0.821	**0.730**	**0.912**
AP	**1.000**	0.998	0.675	0.864	0.922	0.829	0.606	0.780
Aug-AP	**1.000**	**1.000**	0.613	0.799	**0.982**	**0.957**	0.454	0.636
Diversified-AP	0.994	0.970	**0.784**	**0.943**	0.942	0.836	**0.741**	**0.922**
Adv	**1.000**	**0.992**	0.768	0.936	**1.000**	**0.990**	0.459	0.604
Diversified-Adv	0.996	0.976	**0.807**	**0.952**	0.988	0.974	**0.732**	**0.916**

such context-rich inputs are still mapped to a blank or isolated anchor c_a, the erasure loss may treat part of the surrounding context as target-related content. This can create an overly aggressive erasure signal and reduce preservation. As in Table 2, one-sided prompt augmentation (Aug-ESD/AP) and adversarial perturbation (AdvUnlearn) [47] both achieve strong erasure at the expense of preservation, causing *noticeable PSR drops* under multi-object settings.

4. Diversified Unlearning

To overcome the fragility of context-free semantic formulations, we propose **Diversified Unlearning**. Instead of viewing a concept through an isolated keyword, our approach operationalizes that a concept should be unlearned through the company it keeps: erasure is modeled as a structured transformation over diverse contextual variations. By pairing target concepts with natural contextual expressions, we introduce an *implicit contrastive mechanism* that encourages the model to distinguish between the core identity to be erased and the surrounding context to be preserved. Here, *contrastive* refers to paired context-invariant alignment rather than an additional contrastive-loss objective.

Diversified Prompting. To resolve the limitations of keyword-based erasure, we propose **Diversified Prompting**, which replaces the unconstrained keyword objective with a context-invariant formulation. For each target concept c_e, we sample natural contexts $c \sim \mathbf{C}$ to form composite pairs (e.g., pairing c_e "Henry Cavill" with c "waving" to yield "Henry Cavill waving"). The output-space erasure loss is optimized over these contextualized formulations:

$$\min_{\theta'} \mathbb{E}_{c_e \in \mathbf{E}, c_p \in \mathbf{P}, c \in \mathbf{C}} \Bigg[\underbrace{\|\epsilon_{\theta'}(\tau(c+c_e)) - \epsilon_\theta(\tau(c+c_a))\|_2^2}_{L_1} + \delta \cdot \underbrace{\|\epsilon_{\theta'}(\tau(c_p)) - \epsilon_\theta(\tau(c_p))\|_2^2}_{L_2} \Bigg], \tag{4}$$

where $\delta \in \{0,1\}$ activates preservation loss L_2 only if required by the baseline; z_t and t are omitted for simplicity.

This architecture is designed to address the two limitations outlined in our empirical analysis. On the one hand, by keeping the surrounding context c invariant across both sides of L_1, the optimization provides a structural alignment that helps *distinguish what to erase from what to preserve*; it encourages target removal while reducing unintended changes to co-occurring background content, mitigating the output-space ambiguity described in **Observation 2**. On the other hand, embedding c_e within natural phrases *aligns with the text encoder's pre-training distribution*, exposing the optimization to context-sensitive variations that can improve robustness to prompt rephrasing jailbreaks (**Observation 1**).

Diversified Embedding Mixup. While Diversified Prompting works well in the **output-based setting**, applying the same strategy directly to **attention-based unlearning** often causes *severe over-forgetting*. This stems from the **semantic bias of the text encoder**: even non-sensible prompts (e.g., an empty string "") are mapped to meaningful embeddings that yield coherent outputs. In output-based unlearning (Equation (2)), the loss is *bounded by a realistic output distribution*, preventing collapse. However, in attention-based unlearning (Equation (3)), *no such regularization exists*, often leading to *non-sensible generations* for erased concepts, a phenomenon also noted by [11].

To address this, we introduce **diversification directly in the embedding space**. Specifically, we define a mixup function $f(c_e, \mathbf{C})$ that interpolates the token embedding of c_e with *contextualized embeddings* c. The diversified attention-based loss is:

$$\min_{W'} \sum_{c_e \in \mathbf{E}} \left\| W' f(\tau(c_e), \mathbf{C}) - W\tau(c_a) \right\|_2^2 + \sum_{c_p \in \mathbf{P} \cup \mathbf{C}} \left\| W' \tau(c_p) - W\tau(c_p) \right\|_2^2, \tag{5}$$

where the **mixup function** is defined *token-wise*. Let $\mathcal{I}(c_e)$ denote the token positions corresponding to the target concept in the tokenized input. In the mixed branch, $\tau(c)^i$ denotes the contextual-concept embedding aligned to the corresponding target-token slot:

$$f(\tau(c_e), \mathbf{C})^i = \begin{cases} \tau(c_e)^i, & \text{if } i \notin \mathcal{I}(c_e) \\ \dfrac{1}{|\mathbf{C}|} \displaystyle\sum_{c \in \mathbf{C}} \Big(\alpha \tau(c_e)^i + (1-\alpha)\tau(c)^i \Big), & \text{if } i \in \mathcal{I}(c_e). \end{cases} \tag{6}$$

For example, if the input prompt after tokenization is represented as "⟨SOS⟩, ⟨a⟩, ⟨photo⟩, ⟨of⟩, ⟨Henry⟩, ⟨Cavill⟩, ⟨PAD⟩, ..." where the tokens ⟨Henry⟩ and ⟨Cavill⟩ correspond to the **concept** c_e and the remaining tokens are treated as *template tokens*, we replace the embeddings of c_e with mixed features using *contextual concepts* $c \in \mathbf{C}$ (e.g., "a man", "an actor"). Unlike the simple *prompt-level mixup* $\alpha\,\tau(c_e) + (1-\alpha)\,\tau(c)$, this **token-level formulation** yields better performance.

This design *preserves the identity of the target concept* while injecting *contextual diversity*. In practice, we set α close to 1, retaining most of the target signal while adding just enough variation to *mitigate over-forgetting*. As a lightweight add-on, Diversified Embedding Mixup applies preservation regularization over $\mathbf{P}$ only when present in the baseline objective. Since contextual concepts $c \in \mathbf{C}$ are mixed with the target concept c_e, we also *preserve* c to prevent unintended forgetting, which motivates careful selection of contextual concepts.

Context Set Construction. To implement Diversified Unlearning, we construct a context set $\mathbf{C}$ for each target concept c_e using advanced LLMs (e.g., GPT-4o/GPT-5), aiming to approximate how the concept **naturally appears in real-world usage**. We employ a **unified prompt-design rule** and provide a single illustrative example per setting to prompt the LLMs (see Supplementary Material B).

Applied exclusively to training prompts, our filtering pipeline uses a task-aligned recognition model on synthesized images. Our generation rules guarantee high-confidence detections for every candidate, bypassing the need for hyperparameter tuning. For celebrity erasure, we expand simple templates with everyday actions (e.g., "a photo of Henry Cavill smiling"); for artistic styles, we vary the subject while anchoring the style descriptor (e.g., "a art of a fox in the style of Van Gogh"); and for explicit concepts, we blend the keyword with diverse subjects (e.g., "nudity portrait"). This principled construction prevents optimization from overfitting to a fixed keyword token, directly enhancing *robustness and generality* of erasure.

5. Experiments

In this section, we present a **comprehensive evaluation** of our method's ability to enhance existing unlearning approaches, effectively *erasing diverse concepts* while *preserving essential knowledge*. We benchmark our Diversified Unlearning against five recent and representative erasure techniques: ESD [10], UCE [11], AP [4], AGE [5], and ACE [40] under identical experimental settings. In total, we evaluate **10 unlearning methods** across **five target scenarios**, evaluating *erasure fidelity*, *benign concept preservation*, and *robustness against recovery attacks*. Across 50 erasure/preservation comparisons, diversified variants achieve **37 wins, 6 ties, and 7 losses**; the complete per-setting breakdown and W/T/L assignment rule are provided in Table 1 of the Supplementary Material B.1. The empirical results demonstrate that our framework yields systematic advantages over keyword-based counterparts across

WACV #99

Table 3. Quantitative comparison between keyword-based methods and their diversified counterparts for simultaneous erasure of ten celebrities.

Method	Erasure		Preservation			
	CLIP-i↓	GPT↓	LPIPS↓	CLIP-i↑	CLIP-t↑	GPT↑
ESD	65.98	11.78	0.66	62.59	24.25	31.68
Diversified-ESD	**62.58**	**5.50**	**0.63**	**69.73**	**26.28**	**46.15**
UCE	60.59	4.48	0.65	62.90	23.97	**31.42**
Diversified-UCE	**58.55**	**2.83**	**0.64**	**66.92**	**24.57**	24.40
AP	61.14	13.55	0.79	49.31	21.34	14.97
Diversified-AP	**60.78**	**6.10**	**0.68**	**70.50**	**27.09**	**51.15**
AGE	60.68	18.84	0.83	42.31	17.96	6.98
Diversified-AGE	**58.84**	**12.59**	**0.71**	**64.20**	**26.15**	**12.59**
ACE	74.69	29.45	0.46	84.64	**29.92**	**76.72**
Diversified-ACE	**68.67**	**13.53**	**0.39**	**85.91**	28.88	75.05

Table 4. Quantitative results comparing keyword-based methods with their diversified counterparts on erasure of the copyrighted "Mario" character.

Method	Erasure		Preservation			
	CLIP-i↓	GPT↓	LPIPS↓	CLIP-i↑	CLIP-t↑	GPT↑
ESD	66.38	16.25	0.34	91.66	31.54	72.18
Diversified-ESD	**66.29**	**13.50**	**0.32**	**92.38**	**31.61**	**73.65**
UCE	70.23	**21.78**	0.43	93.92	**31.95**	94.78
Diversified-UCE	**71.97**	26.58	**0.37**	**94.98**	31.93	**94.83**
AP	64.89	12.70	0.33	92.50	31.54	73.95
Diversified-AP	**63.42**	**8.60**	**0.31**	**92.56**	**31.73**	**74.35**
AGE	62.49	11.03	0.35	90.62	31.53	70.83
Diversified-AGE	**60.29**	**6.00**	**0.34**	**90.87**	**31.53**	**71.05**
ACE	64.52	**0.28**	0.30	91.50	31.84	89.15
Diversified-ACE	**63.55**	1.05	**0.27**	**93.12**	**32.02**	**91.23**

these three evaluation aspects. Due to page limits, quantitative analysis for three representative settings—*celebrity*, *copyrighted*, and *nudity* concepts—is detailed below, while the remaining results are provided in the Supplementary Material (Supp. Material).

5.1. Celebrity and Copyrighted Character Erasure

Experimental Setup. We evaluate identity unlearning across two scenarios: multi-concept celebrity and copyrighted character erasure. For celebrities, we simultaneously target ten distinct figures (single-concept details in the Supp. Material B.1.1), assessing erasure on a 2,000-image benchmark (eight complexity levels) and preservation across 1,500 prompts (15 un-erased subjects). For copyrighted content, we target "Mario", evaluating erasure on 1,000-image (five complexity levels) and preservation across 1,000 prompts (10 characters).

Across both domains, output-based variants [5] (Diversified-ESD, -AP, -AGE, -ACE) utilize 20 fine-tuning prompt sets per concept, while Diversified-UCE uses 5. Given their task similarity, both setups share identical evaluation metrics: LPIPS [46] and CLIP [28] for visual fidelity, alongside an identity-focused GPT-Score from [25] implemented via Qwen2.5-VL-72B [2] to rate resemblance on a 0–4 scale normalized to percentages. The best results are highlighted in **bold** within the respective tables.

Results. Empirical outcomes on celebrity erasure (Table 3) show that our diversified framework broadens performance across both erasure and preservation. Notably, ACE exhibits a dramatic reduction in GPT-Score from 29.45% to 13.53%, highlighting the effectiveness of diversification. While UCE achieves the strongest erasure among the baselines, its diversified variant delivers an even more pronounced effect, yielding the lowest GPT-Score. When extended to adversarial-training baselines, our approach maintains erasure efficacy while improving preservation (see Table 2, Supp. Material for detailed results). Correspondingly, metrics for "Mario" (Table 4) confirm that Diversified Unlearning advances copyrighted erasure; Diversified-AP and -AGE nearly **halve the identity GPT-scores** relative to their baselines (from 12.70 to 8.60, and 11.03 to 6.00). This generalizes to trademark unbranding [23] (Section C.4, Supp. Material), where our method decouples logos from physical products, boosting overall unbranding success.

Qualitative comparisons (Figure 2, top row), (Figure 3, and Figure 5, Supp. Material) mirror these quantitative trends. While target identities remain recognizable under baseline ESD [10], AP [4], and AGE [5], they become unrecognizable under our framework. Interestingly, although the specific character identity of "Mario" is removed, our approach still safeguards peripheral stylistic primitives such as colors and textures. This demonstrates a *smooth erasure effect* rather than the overly aggressive, destructive "hard wash" typical of keyword-based baselines, directly benefiting the retention of adjacent benign concepts.

5.2. Explicit Content Erasure

Experimental Setup. For Not-Safe-For-Work (NSFW) removal ("nudity"), we fine-tune cross-attention layers in Stable Diffusion [31] with 20 prompt sets (see Supp. Material B.1.5). Erasure efficacy is evaluated on the I2P benchmark [33] (4,703 images) and preservation on MS-COCO 30K [20]. We measure nudity occurrences via NudeNet [27], image distribution similarity via FID [15], and semantic caption alignment via CLIP score [28].

Results. As detailed in Table 6, our diversified framework suppresses explicit content while anchoring generation quality on COCO-30K. Despite a slight CLIP regression in Diversified-UCE (29.30 vs. 31.27), diversified methods maintain strong FID/CLIP scores comparable to the original SD [31] backbone. Quantitatively, Diversified-ESD and -AP collapse total nudity detections by 44% and

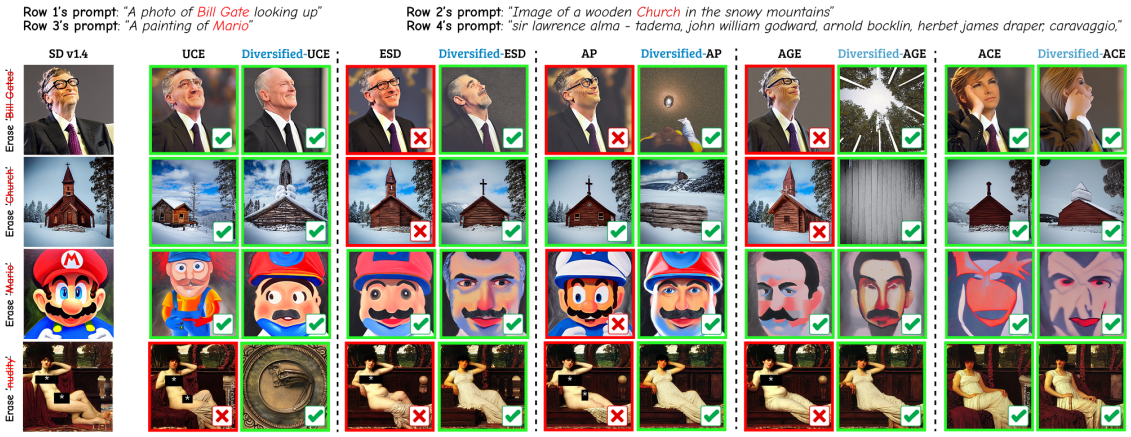

Figure 2. Visualization outcomes of Diversified Unlearning compared to keyword-based methods across four settings: Celebrities, Objects, Copyrighted Characters, and Explicit Content erasure.

Table 5. Robustness evaluation under recovery attacks. We report Attack Success Rate for Ring-A-Bell [39], and GPT-score [25] for both Indirect Recovery (IR) of the "Mario" character and Noise-based attacks (Noise ATK) on ten-celebrity erasure.

Method	Ring-A-Bell			IR
	K16↓	K38↓	K77↓	GPT↓
AP	46.32	48.42	54.47	38.13
Diversified-AP	**38.95**	**41.05**	**38.95**	**33.75**

Noise ATK (GPT ↓)	**AP**	**ESD**	**AGE**	**UCE**
Baseline	39.63	28.19	33.00	2.00
Diversified	**14.94**	**11.75**	**29.81**	**1.13**

Table 6. Quantitative results of exposed body-part detection comparing keyword-based methods with their diversified counterparts. Erasure is evaluated using Nudenet [27] on the I2P dataset [33], where fewer detections indicate better performance, while preservation is assessed on COCO-30K [20].

Method	Erasure of nudity with NudeNet on I2P									MS-COCO 30K	
	Armpits	Belly	Buttocks	Feet	Breasts(F)	Genitalia(F)	Breasts(M)	Genitalia(M)	Total↓	FID↓	CLIP↑
SD	148	170	29	63	266	18	42	7	743	14.04	31.34
UCE	69	61	7	23	**56**	6	**10**	21	253	**14.85**	**31.27**
Diversified-UCE	**45**	**45**	**6**	**6**	62	**1**	22	**9**	**196**	14.87	29.30
ESD	105	70	16	24	128	6	15	13	377	14.71	30.82
Diversified-ESD	**60**	**38**	**7**	**16**	**72**	**3**	**10**	**6**	**212**	**14.05**	**31.06**
AP	91	61	19	26	123	10	12	**13**	355	14.71	30.98
Diversified-AP	**52**	**30**	**5**	**10**	**49**	**3**	**8**	15	**172**	**14.21**	**31.01**
AGE	84	50	14	21	101	7	14	**9**	300	14.41	30.99
Diversified-AGE	**68**	**40**	**9**	**17**	**77**	**5**	**10**	18	**244**	**13.82**	**31.00**
ACE	7	7	4	**1**	12	3	**0**	**5**	39	**14.63**	**30.85**
Diversified-ACE	**5**	**4**	**2**	4	**5**	**0**	**0**	6	**26**	14.85	30.60

52% relative to their respective baselines. For the highly optimized ACE architecture [40], Diversified-ACE further improves the baseline's erasure performance, reducing the target concept presence score from 39 to 26, while maintaining a comparable preservation profile. Crucially, this performance parity scales to robustness defenses, where Diversified-ACE matches the baseline's resistance against recovery attacks (Table 9, Supp. Material). Visualizations (Figure 1, Figure 6 in the Supp. Material) visually confirm these metrics, demonstrating clear erasure of target concepts alongside clean retention of unrelated contexts.

5.3. Robustness against Recovery Attacks

Recent studies [21] categorize concept unlearning into *guidance-based avoidance* (steering generation away) and *destruction-based removal* (directly decaying target likelihood). While keyword-based methods often inadvertently collapse into surface-level avoidance–leaving latent concepts vulnerable to prompt-rephrasing or recovery jailbreaks [14, 26, 39]–we hypothesize that *Diversified Unlearning* functions as a robust destruction-based mechanism. By unlearning a concept across its contextual company rather than through a single fixed token, our framework enforces concept removal across diverse linguistic variations, thereby establishing tighter erasure security.

Figure 5. Intermediate generations comparing AP and Diversified-AP after erasing "Margot Robbie," under both normal prompting and Noise-based attacks.

Experimental Setup. We evaluate robustness under three distinct attack vectors: (1) **Adversarial prompts** via Ring-A-Bell [39] for the "nudity"; (2) **Indirect recovery**

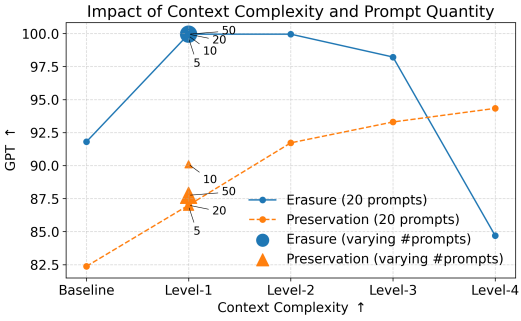

Figure 3. Effect of context complexity and prompt quantity on ESD-based unlearning. Context complexity increases from Level-1 to Level-4 with 20 prompts, while prompt quantity varies from 5 to 50 at Level-1.

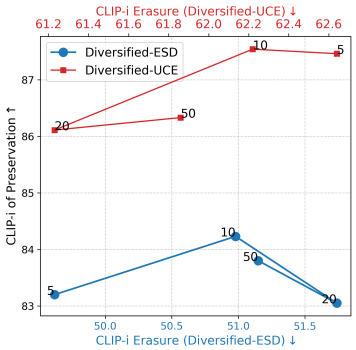

Figure 4. Impact of prompt set size on Diversified-ESD at context Level-1 and Diversified-UCE ($\alpha = 0.999$) for "Henry Cavill" erasure.

Table 7. Impact of different α values in Diversified Embedding Mixup on Diversified-UCE under the ten-celebrity erasure setting. Overall, decreasing α yields a marginal gain in erasure but causes reduced preservation. Setting $\alpha = 0.999$ establishes an optimal trade-off.

α	Erasure		Preservation			
	CLIP-i↓	GPT↓	LPIPS↓	CLIP-i↑	CLIP-t↑	GPT↑
UCE	60.59	4.48	0.65	62.90	23.97	**31.42**
0.999	58.55	2.83	**0.64**	**66.92**	**24.57**	24.40
0.990	58.01	2.79	**0.64**	64.60	23.95	23.00
0.950	**57.25**	3.05	0.65	60.58	22.32	15.37
0.900	57.66	2.53	0.66	60.09	21.88	13.10
0.800	57.45	**2.46**	0.67	57.84	20.86	10.07

(IR) [14] via high-context prompts that omit the name "Mario" but still trigger the character; and (3) **Noise-based attacks** [21] on the ten-celebrity erasure models. Defense is measured by Attack Success Rate (ASR) for Ring-A-Bell, and identity-focused GPT-Score [25] for IR and Noise-based attacks (lower is better).

Results. As shown in Table 5, Diversified-AP **consistently outperforms** its baseline across all three targeted domains (nudity, copyrighted characters, and celebrities), demonstrating resilience against diverse recovery strategies.

To better understand the underlying mechanism, we compare AP and our Diversified-AP, which unlearn the same "Margot Robbie" concept, under the same input prompt and with and without Noise-based attacks (Figure 5). While the final outputs of both methods successfully obscure the target concept under normal conditions (rows 1–2), the intermediate generation steps (rows 3–4) expose a critical disparity: standard AP can still be jailbroken to reconstruct the target latent features, signifying non-robust unlearning. This trajectory aligns with our quantitative data: under Noise-based attacks, Diversified-AP suppresses concept recovery far more effectively, yielding a lower recovery score (14.94%) than the baseline AP (39.63%). Our framework thus tightens structural robustness alongside traditional erasure and preservation metrics. Additional quantitative validations are detailed in the Supp. Material C.6.

5.4. Ablation Study

Experimental Setup. We evaluate three core hyperparameters: **contextual diversity**, **mixup strength**, and **prompt quantity**. For contextual diversity, we test Diversified-ESD on "Henry Cavill," comparing the baseline [10] (canonical prompt: "A photo of Henry Cavill") against four variants fine-tuned with prompts of increasing complexity (Supp. Material B.2). For embedding mixup, we vary the interpolation coefficient α in Diversified-UCE under a simultaneous ten-celebrity erasure setting. Finally, we examine prompt quantity sensitivity by fixing complexity at Level-1 for "Henry Cavill" (or setting $\alpha = 0.999$) and varying the number of fine-tuning prompts for Diversified-ESD and -UCE following the protocols in Section 5.1.

Results. Empirical data (Figure 3) show that diversified prompting yields immediate advantages over the keyword-only baseline even at Level-1, validating the utility of contextual variation. Scaling prompt complexity up to a moderate tier (Level-3) provides **the optimal erasure–preservation trade-off**, whereas excessive textual complexity (Level-4) degrades erasure performance. Furthermore, outcomes in Table 7 indicate that decreasing α amplifies contextual entanglement with the target concept but severely **harms benign space preservation**. Because erasure improves marginally, gains quickly saturate around $\alpha = 0.95$, we anchor $\alpha = 0.999$ across all main experiments. Finally, Figure 4 demonstrates that performance remains **highly stable across a wide range of prompt quantities** provided the context complexity is properly bounded.

6. Conclusion

In this work, we revisit the brittleness of keyword-based concept unlearning in text-to-image models, highlighting how fixed-indicator formulations fail to address context-dependent embedding regions and output-space entanglement. Guided by the principle that a concept should be unlearned through the company it keeps, we introduce *Diversified Unlearning*, a lightweight framework that injects implicit contrastive supervision into the unlearning pipeline via structured context pairs. Implemented through two unified, plug-and-play modules enhancing existing techniques, Diversified Prompting and Diversified Embedding Mixup, our approach establishes a sharper optimization boundary within both output and attention spaces. Extensive evaluations across five settings confirm that our framework broadly delivers the triple benefits of stronger erasure, cleaner retention of benign neighbors, and improved robustness against advanced adversarial jailbreaks without architectural overhead.

WACV #99

References

[1] Youssef Allouah, Rachid Guerraoui, and Sanmi Koyejo. Distributional unlearning: Forgetting distributions, not just samples. In *ICML 2025 Workshop on Machine Unlearning for Generative AI*, 2025. 2, 4

[2] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibo Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. Qwen2.5-vl technical report. *arXiv preprint arXiv:2502.13923*, 2025. 6

[3] Wieland Brendel, Jonas Rauber, and Matthias Bethge. Decision-based adversarial attacks: Reliable attacks against black-box machine learning models. *arXiv preprint arXiv:1712.04248*, 2017. 3

[4] Anh Bui, Long Vuong, Khanh Doan, Trung Le, Paul Montague, Tamas Abraham, and Dinh Phung. Erasing undesirable concepts in diffusion models with adversarial preservation. *NeurIPS*, 2024. 1, 3, 4, 5, 6

[5] Anh Bui, Trang Vu, Long Vuong, Trung Le, Paul Montague, Tamas Abraham, and Dinh Phung. Fantastic targets for concept erasure in diffusion models and where to find them. *ICLR*, 2025. 1, 3, 5, 6

[6] Nicholas Carlini, Jamie Hayes, Milad Nasr, Matthew Jagielski, Vikash Sehwag, Florian Tramèr, Borja Balle, Daphne Ippolito, and Eric Wallace. Extracting training data from diffusion models. In *Proceedings of the 32nd USENIX Security Symposium (USENIX Security '23)*, pages 5253–5270, 2023. Full paper (PDF): https://www.usenix.org/system/files/usenixsecurity23-carlini.pdf. 1

[7] Zhi-Yi Chin, Chieh-Ming Jiang, Ching-Chun Huang, Pin-Yu Chen, and Wei-Chen Chiu. Prompting4debugging: Red-teaming text-to-image diffusion models by finding problematic prompts. In *International Conference on Machine Learning (ICML)*, 2024. 3

[8] Chongyu Fan, Jiancheng Liu, Yihua Zhang, Eric Wong, Dennis Wei, and Sijia Liu. Salun: Empowering machine unlearning via gradient-based weight saliency in both image classification and generation. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024. Spotlight. 3

[9] J. R. Firth. A synopsis of linguistic theory, 1930–1955. In *Studies in Linguistic Analysis*, pages 1–32. Blackwell, Oxford, 1957. 1

[10] Rohit Gandikota, Joanna Materzyńska, Jaden Fiotto-Kaufman, and David Bau. Erasing concepts from diffusion models. In *Proceedings of the 2023 IEEE International Conference on Computer Vision*, 2023. 1, 3, 4, 5, 6, 8

[11] Rohit Gandikota, Hadas Orgad, Yonatan Belinkov, Joanna Materzyńska, and David Bau. Unified concept editing in diffusion models. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 5111–5120, 2024. 1, 3, 5

[12] Chao Gong, Kai Chen, Zhipeng Wei, Jingjing Chen, and Yu-Gang Jiang. Reliable and efficient concept erasure of text-to-image diffusion models. In *Proceedings of the European conference on computer vision (ECCV)*, pages 73–88, 2024. 3

[13] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 770–778, 2016. 4

[14] Luxi He, Yangsibo Huang, Weijia Shi, Tinghao Xie, Haotian Liu, Yue Wang, Luke Zettlemoyer, Chiyuan Zhang, Danqi Chen, and Peter Henderson. Fantastic copyrighted beasts and how (not) to generate them. In *2nd Workshop on Generative AI and Law (GenLaw), 41st International Conference on Machine Learning (ICML)*, 2024. 3, 7, 8

[15] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. Gans trained by a two time-scale update rule converge to a local nash equilibrium. In *Advances in Neural Information Processing Systems*, 2017. 6

[16] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020. 3

[17] Chi-Pin Huang, Kai-Po Chang, Chung-Ting Tsai, Yung-Hsuan Lai, Fu-En Yang, and Yu-Chiang Frank Wang. Receler: Reliable concept erasing of text-to-image diffusion models via lightweight erasers. In *Proceedings of the European conference on computer vision (ECCV)*, pages 360–376, 2024. 2, 3, 4

[18] Changhoon Kim, Kyle Min, and Yezhou Yang. R.a.c.e.: Robust adversarial concept erasure for secure text-to-image diffusion model. In *Proceedings of the European conference on computer vision (ECCV)*, pages 461–478, 2024. 3

[19] Nupur Kumari, Bingliang Zhang, Sheng-Yu Wang, Eli Shechtman, Richard Zhang, and Jun-Yan Zhu. Ablating concepts in text-to-image diffusion models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 22691–22702, 2023. 3

[20] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer Vision–ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6-12, 2014, Proceedings, Part V 13*, pages 740–755. Springer, 2014. 6, 7

[21] Kevin Lu, Nicky Kriplani, Rohit Gandikota, Minh Pham, David Bau, Chinmay Hegde, and Niv Cohen. When are concepts erased from diffusion models? *arXiv preprint arXiv:2505.17013*, 2025. 3, 4, 7, 8

[22] Shilin Lu, Zilan Wang, Leyang Li, Yanzhu Liu, and Adams Wai-Kin Kong. Mace: Mass concept erasure in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6430–6440, 2024. 1, 3

[23] Dawid Malarz, Artur Kasymov, Filip Manjak, Maciej Zięba, and Przemysław Spurek. From unlearning to unbranding: A benchmark for trademark-safe text-to-image generation, 2025. 6

WACV #99

[24] Hadas Orgad, Bahjat Kawar, and Yonatan Belinkov. Editing implicit assumptions in text-to-image diffusion models. In *IEEE International Conference on Computer Vision, ICCV 2023, Paris, France, October 1-6, 2023*, pages 7030–7038. IEEE, 2023. 3

[25] Yuang Peng, Yuxin Cui, Haomiao Tang, Zekun Qi, Runpei Dong, Jing Bai, Chunrui Han, Zheng Ge, Xiangyu Zhang, and Shu-Tao Xia. Dreambench++: A human-aligned benchmark for personalized image generation. In *The Thirteenth International Conference on Learning Representations*, 2025. 2, 4, 6, 7, 8

[26] Minh Pham, Kelly O Marshall, and Chinmay Hegde. Circumventing concept erasure methods for text-to-image generative models. *arXiv preprint arXiv:2308.01508*, 2023. 7

[27] Bedapudi Praneet. Nudenet: Neural nets for nudity classification, detection and selective censorin. 2019. 3, 6, 7

[28] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, pages 8748–8763. PMLR, 2021. 3, 6

[29] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-shot text-to-image generation. In *International Conference on Machine Learning*, pages 8821–8831. PMLR, 2021. 1, 3

[30] Aditya Ramesh, Prafulla Dhariwal, Alex Nichol, Casey Chu, and Mark Chen. Hierarchical text-conditional image generation with clip latents. *arXiv preprint arXiv:2204.06125*, 1 (2):3, 2022.

[31] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695, 2022. 1, 3, 6

[32] Chen Ruidong, Guo Honglin, Wang Lanjun, Zhang Chenyu, Nie Weizh, and Liu An-An. Trce: Towards reliable malicious concept erasure in text-to-image diffusion models. *arXiv preprint arXiv:2503.07389*, 2025. 3

[33] Patrick Schramowski, Manuel Brack, Björn Deiseroth, and Kristian Kersting. Safe Latent Diffusion: Mitigating Inappropriate Degeneration in Diffusion Models. In *CVPR*, pages 22522–22531, 2023. 1, 3, 6, 7

[34] Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade W Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton Mullis, Mitchell Wortsman, Patrick Schramowski, Srivatsa R Kundurthy, Katherine Crowson, Ludwig Schmidt, Robert Kaczmarczyk, and Jenia Jitsev. LAION-5b: An open large-scale dataset for training next generation image-text models. In *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2022. 3

[35] Zhan Shi, Xu Zhou, Xipeng Qiu, and Xiaodan Zhu. Improving image captioning with better use of caption. pages 7454–7464, Online, 2020. Association for Computational Linguistics. 3

[36] SmithMano. Tutorial: How to remove the safety filter in 5 seconds, 2022. 3

[37] Gowthami Somepalli, Vasu Singla, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Diffusion art or digital forgery? investigating data replication in diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6048–6058, 2023. 1

[38] StabilityAI. Stable diffusion 2.0 release. 2022. 3

[39] Yu-Lin Tsai*, Chia-Yi Hsu*, Chulin Xie, Chih-Hsun Lin, Jia-You Chen, Bo Li, Pin-Yu Chen, Chia-Mu Yu, and Chun-Ying Huang. Ring-a-bell! how reliable are concept removal methods for diffusion models? In *The Twelfth International Conference on Learning Representations*, 2024. 3, 4, 7

[40] Zihao Wang, Yuxiang Wei, Fan Li, Renjing Pei, Hang Xu, and Wangmeng Zuo. Ace: Anti-editing concept erasure in text-to-image models. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, pages 23505–23515, 2025. 1, 3, 5, 7

[41] Yongliang Wu, Shiji Zhou, Mingzhuo Yang, Lianzhe Wang, Heng Chang, Wenbo Zhu, Xinting Hu, Xiao Zhou, and Xu Yang. Unlearning concepts in diffusion model via concept domain correction and concept preserving gradient. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025. 3

[42] Yijun Yang, Ruiyuan Gao, Xiaosen Wang, Tsung-Yi Ho, Nan Xu, and Qiang Xu. MMA-Diffusion: MultiModal Attack on Diffusion Models. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 3

[43] Yuchen Yang, Bo Hui, Haolin Yuan, Neil Gong, and Yinzhi Cao. Sneakyprompt: Jailbreaking text-to-image generative models. In *2024 IEEE Symposium on Security and Privacy (SP)*, pages 123–123. IEEE Computer Society, 2024. 3

[44] Jaehong Yoon, Shoubin Yu, Vaidehi Patil, Huaxiu Yao, and Mohit Bansal. Safree: Training-free and adaptive guard for safe text-to-image and video generation. *ICLR*, 2025. 3

[45] Gong Zhang, Kai Wang, Xingqian Xu, Zhangyang Wang, and Humphrey Shi. Forget-me-not: Learning to forget in text-to-image diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, pages 1755–1764, 2024. 3

[46] Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 586–595, 2018. 6

[47] Yimeng Zhang, Xin Chen, Jinghan Jia, Yihua Zhang, Chongyu Fan, Jiancheng Liu, Mingyi Hong, Ke Ding, and Sijia Liu. Defensive unlearning with adversarial training for robust concept erasure in diffusion models. In *Advances in Neural Information Processing Systems*, 2024. 2, 3, 4

[48] Yimeng Zhang, Jinghan Jia, Xin Chen, Aochuan Chen, Yihua Zhang, Jiancheng Liu, Ke Ding, and Sijia Liu. To generate or not? safety-driven unlearned diffusion models are still easy to generate unsafe images... for now. *European Conference on Computer Vision (ECCV)*, 2024. 3

[49] Hongguang Zhu, Yunchao Wei, Mengyu Wang, Siyu Jiao, Yan Fang, Jiannan Huang, and Yao Zhao. Sage: Exploring the boundaries of unsafe concept domain with semantic-augment erasing. *arXiv preprint arXiv:2506.09363*, 2025. 3